

Clarification on Data Sources for Federated Learning (FL) in ChainFLIP

This document provides a detailed explanation of where the raw data for training Federated Learning (FL) models in the ChainFLIP project would originate. It addresses your question about block explorer logs (like Amoy Polygon Scan) and outlines the broader data landscape for FL clients.

1. Overview of Data Sourcing for FL

Federated Learning operates on the principle that raw data stays local to each participating client (e.g., Manufacturer, Transporter, Retailer). Clients train local models on their data and only share model updates (like gradients or weights) with a central aggregator. The aggregator combines these updates to create a global model, which is then sent back to clients.

Therefore, the "raw data" is not collected into a central repository. Instead, each FL client application, running within each participating organization's infrastructure, is responsible for accessing and processing its own relevant data to create features for its local ML model.

2. Primary Data Sources for FL Clients

FL clients in the ChainFLIP system would leverage a combination of on-chain data (publicly verifiable) and potentially local, private off-chain data (processed locally).

2.1. On-Chain Data (Accessible via Blockchain Interaction)

This is data recorded on the Polygon Amoy blockchain and is publicly viewable through block explorers like Polygon Scan (Amoy Testnet). However, FL clients would typically access this data programmatically by interacting with a Polygon RPC node (e.g., using libraries like `ethers.js` if the client is Node.js-based, or `web3.py` if Python-based).

How it's used:

- **FL clients query the blockchain:** They would subscribe to events or query historical data from the deployed ChainFLIP smart contracts (`SupplyChainNFT.sol`

and its inherited contracts like `NodeManagement.sol` , `BatchProcessing.sol` , `Marketplace.sol` , `DisputeResolution.sol`).

- **Data types from smart contracts:**

- **Event Logs:** This is a primary source. For example:

- From `NodeManagement.sol` : Events related to node registration (if an event is emitted), role changes, type changes, verification status updates, reputation updates (`NodeReputationUpdated` , `NodePenalized`).
 - From `NFTCore.sol` / `Marketplace.sol` : `ProductMinted` , `Transfer` (standard ERC721), `PurchaseInitiated` , `TransportStarted` , `TransportCompleted` , `CIDStored` , `ProductHistoryCIDUpdated` .
 - From `BatchProcessing.sol` : `BatchProposed` , `BatchValidated` (validator, vote, batchId), `BatchCommitted` , `BatchFlagged` .
 - From `DisputeResolution.sol` : `DisputeOpened` , `ArbitratorProposed` , `ArbitratorVoted` , `DisputeResolved` .

- **State Variables:** Clients can also query public state variables of the contracts:

- `NodeManagement.sol` : Current reputation of a node (`nodeReputation(address)`), node roles (`nodeRoles(address)`), node types (`nodeTypes(address)`), verification status (`isVerifiedNode(address)`).
 - `NFTCore.sol` : Owner of an NFT (`ownerOf(tokenId)`), product details linked to an NFT (if stored on-chain, though most are on IPFS via CID).
 - `BatchProcessing.sol` : Details of a batch (`getBatchDetails(batchId)`).

- **Relevance to Block Explorers:** While a block explorer shows this data, the FL client interacts directly with the blockchain network. The explorer is a user interface for the same underlying data.

Example: Sybil Detection Feature Engineering from On-Chain Data

A client might gather data for a newly registered node: 1. **Registration Time:** If `NodeManagement.sol` emits an event upon new node registration or has a public mapping storing registration timestamps (this might need to be added, as noted in the guide). 2. **Initial Transactions:** Monitor `Transfer` , `BatchProposed` , `ProductMinted` events involving this new node in its first few days/weeks. 3. **Features Created Locally:** `registration_age` , `transaction_count_first_week` , `interaction_diversity_first_week` (how many different types of actions it performed), `early_batch_proposal_rate` .

2.2. Off-Chain Data (Local to Each Participant)

This data resides within each participating organization's own systems and is **not directly shared** with other participants or the FL aggregator.

- **Types of Local Data:**

- **Internal Transaction Logs:** More detailed logs of their own supply chain activities beyond what's on-chain.
- **Communication Records (Metadata):** Patterns of communication with other supply chain partners (e.g., frequency, timing – not content).
- **KYC/Verification Data:** Information collected during their own due diligence when onboarding partners (used very carefully and only for local feature generation, never shared).
- **Internal Risk Assessments:** If an organization has its own internal system for flagging suspicious partners or transactions.
- **Operational Data:** Logistics details, inventory levels, etc., that might correlate with legitimate or fraudulent behavior.

- **How it's used by the FL Client:**

- The FL client application, running within the organization's secure environment, can be given access to these local data sources.
- It processes this data locally to generate features. For example, a sudden spike in off-chain communication with a new, low-reputation node followed by on-chain transactions might be a feature.
- **Crucially, the raw local data itself never leaves the organization's premises.** Only the features derived from it are used to train the local model, and only the model updates are shared.

Example: Bribery Detection Feature Engineering (Hypothetical)

If a Transporter's FL client has access to (anonymized) local payment logs (highly sensitive and requires strict local privacy controls):

1. **Local Data:** Observes an unusual, off-chain payment to a Validator node around the time of a critical batch validation.
2. **On-Chain Data:** Correlates this with the Validator's on-chain vote for that batch.
3. **Feature Created Locally:** `has_proximal_off_chain_payment_to_validator_for_batch_X` (boolean). This is a complex and sensitive feature, illustrating how local data could be used. The focus should always be on privacy-preserving feature engineering.

2.3. Data from `backendListener.js` (Indirectly)

The `backendListener.js` script primarily interacts with IPFS based on on-chain events. It's not a primary raw data source for FL clients in the same way as direct blockchain access or local databases.

- **Potential Use:**

- **Aggregated Statistics (if implemented):** If `backendListener.js` were enhanced to produce anonymized, aggregated statistics about IPFS interactions (e.g., average size of history files, frequency of updates per product type, common access patterns for CIDs), these statistics could be published (e.g., to a secure API or even on-chain if small enough) and consumed by FL clients as contextual features.
- **Logs for Admins:** The logs from `backendListener.js` would be more for administrative monitoring and debugging the IPFS interaction layer rather than direct input for FL clients, unless they indicate specific node misbehavior related to IPFS (e.g., a node consistently failing to provide valid CIDs).

3. The Role of Feature Engineering

It's vital to understand that FL models are not trained on raw blockchain transaction logs or raw database entries directly. Instead, each FL client performs **local feature engineering**:

1. **Data Collection:** Gathers relevant data from the sources mentioned above.
2. **Preprocessing:** Cleans and transforms the data (e.g., handling timestamps, categorical variables).
3. **Feature Extraction:** Creates meaningful numerical features that capture behavioral patterns, anomalies, or characteristics relevant to the FL task (e.g., Sybil detection, collusion detection).
 - Examples were provided in the FL integration guide (e.g., `registration_age_days`, `validator_agreement_rate`).
4. **Local Model Training:** These engineered features are then used to train the local ML model.

4. Privacy and Security

The core design of FL ensures that:

- **Raw data stays local:** Participants do not expose their sensitive operational data or detailed blockchain interaction logs to others.

- **Only model updates are shared:** These updates are aggregated and ideally should not allow reverse-engineering of the underlying raw data from any single client, especially if techniques like differential privacy are added locally or secure aggregation is used.

In summary, FL clients in ChainFLIP will primarily source data by programmatically interacting with the Polygon Amoy blockchain (for on-chain events and state) and by accessing local, private data within their own organizational environments. This data is then processed locally into features to train their respective ML models, contributing to the global FL model in a privacy-preserving manner.